

Intro to NLP

Assignment - 4

Soumodipta Bose

2021201086

Exploring Data and Pre-processing

The first step is cleaning the data, for that i am doing the following :

```
def clean_text(self,line):
    line = line.strip()
    line = re.sub(r'<|>', ' ', line)
    line = self.replace_dates(line)
    line = self.replace_hyphenated_words(line)
    line = self.replace_hash_tags(line)
    # remove < and > from the text
    line = clean(line, no_emoji=True,
                 no_urls=True,
                 no_emails=True,
                 no_phone_numbers=True,
                 no_currency_symbols=True,
                 replace_with_url=" <URL> ",
                 replace_with_email=" <EMAIL> ",
                 replace_with_phone_number=" <PHONE> ",
                 replace_with_currency_symbol=" <CURRENCY> ",
                 lower=True)
    line = self.remove_special_characters(line)
    line = clean(line,no_numbers=True,
                 no_digits=True,
                 no_punct=True,
                 replace_with_number=" <NUMBER> ",
                 replace_with_digit=" ",
                 replace_with_punct="",
                 lower=True)
    line = self.remove_extra_spaces(line)
    tokens=self.tokenizer(line)
    return " ".join(tokens)
```

```

def remove_stopwords(self,text):
    tokens = self.tokenizer(text)
    return " ".join([token for token in tokens if token not in self.stopwords])

def lemmatize(self,text):
    doc = self.nlp(text)
    return " ".join([token.lemma_ for token in doc])

def process(self,text):
    text = self.clean_text(text)
    text = self.remove_stopwords(text)
    text = self.lemmatize(text)
    return text

```

Next i am removing the stop words and lemmatizing the sentences.

```

class DataPipeline(Dataset):
    def __init__(self, filename,type,max_seq_len=50,min_freq=3,vocab=None):
        self.read_data(filename,type)
        self.max_seq_len = max_seq_len
        if vocab is None:
            self.vocab, self.ind2vocab,self.word_count = self.build_vocab(self.data,min_freq)
        else:
            self.vocab = vocab
            self.ind2vocab = {v: k for k, v in vocab.items()}
            # self.word_count = self.get_word_count(vocab,self.data)
        self.ind2vocab = {ind: word for word, ind in self.vocab.items()}

```

Using a common DataPipeline class from which i am creating two Dataset Specific dataset files for example SST data :

```

class SstData(DataPipeline):
    def read_data(self, filename, type):
        datacleaner = DataCleaner()
        data =load_from_disk(filename)
        processed_data = []
        target = []
        for line in tqdm(data[type]):
            processed_data.append(datacleaner.process(line['sentence']).split(" "))
            target.append(line['label'])
        self.data=processed_data
        self.target=target
    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        sent = self.data[idx]
        label = self.target[idx]
        # padding the sentences to create sequences of same length
        if len(sent) < self.max_seq_len:
            sent=[self.word_to_ind(token) for token in sent]+[self.word_to_ind("<pad>") for _ in range(self.max_seq_len-len(sent))]
        return torch.LongTensor(sent),torch.Tensor([label])

```

Since the Elmo Dataset needs to be created for both the dataset i am creating a common interface to convert the dataset to a elmo compatible dataset for training.

```
class ElmoDataset(Dataset):
    def __init__(self, dataset :DataPipeline):
        self.data = dataset.data
        self.max_seq_len = dataset.max_seq_len
        self.vocab = dataset.vocab
        self.ind2vocab = dataset.ind2vocab
        self.word_to_ind = dataset.word_to_ind
        self.ind_to_word = dataset.ind_to_word
    def __len__(self):
        return len(self.data)
    def __getitem__(self, idx):
        sent = self.data[idx]
        # padding the sentences to create sequences of same length
        if len(sent) < self.max_seq_len:
            sent=[self.word_to_ind(token) for token in sent]+[self.word_to_ind("<pad>") for _ in range(self.max_seq_len-len(sent))]
        forward_data = sent[1:]
        backward_data = sent[:-1]
        return torch.LongTensor(forward_data),torch.LongTensor(backward_data)
    def get_batches(self, batch_size):
        return DataLoader(self, batch_size=batch_size, shuffle=False,drop_last=True)
```

Model Creation

The model was created using Pytorch. The model consisted of an initial embedding layer that used pre-trained embedding embeddings using Glove. Then was fed into an LSTM and finally a Linear Fully connected layer with output dimensions of the tag set size.


```
class ELMo(nn.Module):
    def __init__(self, vocab_size, embedding_dim, hidden_dim, max_len, embedding_matrix):
        super(ELMo, self).__init__()
        self.vocab_size = vocab_size
        self.embedding_dim = embedding_dim
        self.hidden_dim = hidden_dim
        self.max_len = max_len
        self.embedding = nn.Embedding.from_pretrained(embedding_matrix)
        self.embedding.weight = nn.Parameter(self.embedding.weight, requires_grad=True)
        self.lstm1 = nn.LSTM(embedding_dim, hidden_dim, batch_first=True, bidirectional=True)
        self.lstm2 = nn.LSTM(hidden_dim*2, hidden_dim, batch_first=True, bidirectional=True)
        self.linear_out = nn.Linear(hidden_dim*2, vocab_size)
    def forward(self, back_data):
        back_embed = self.embedding(back_data)
        back_lstm1, _ = self.lstm1(back_embed)
        #back_lstm1 is shape of (batch_size, max_len, hidden_dim*2)
        back_lstm2, _ = self.lstm2(back_lstm1)
        linear_out = self.linear_out(back_lstm2)
        return linear_out
```

Training the model

```
class ElmoTrainer:
    def __init__(self, epochs=20, lr=0.001, batch_size=50, print_every=1, device='cpu'):
        self.epochs = epochs
        self.lr = lr
        self.batch_size = batch_size
        self.print_every = print_every
        self.device = device
        self.criterion = nn.CrossEntropyLoss()
        self.lowest_validation_loss = float('inf')

    def train(self, model : ELMo, model_save_path, train_data, validation_data):
        self.optimizer = optim.Adam(model.parameters(), lr=self.lr)
        model.to(self.device)
        for epoch in range(len(range(self.epochs))):
            model.train()
            train_loader = train_data.get_batches(self.batch_size)
            training_loss = 0
            for (forward_data, backward_data) in tqdm(train_loader):
                forward_data = forward_data.to(self.device)
                backward_data = backward_data.to(self.device)
                self.optimizer.zero_grad()
                output = model(backward_data)
                output = output.view((-1, model.vocab_size))
                target = forward_data.view(-1)
                loss = self.criterion(output, target)
                loss.backward()
                self.optimizer.step()
                training_loss += loss.item()
            if epoch % self.print_every == 0:
                print('Training Loss : {}'.format(training_loss/len(train_loader)))
            self.__validate(model, model_save_path, validation_data)
```

```
100%|██████████| 170/170 [00:06<00:00, 24.55it/s]
22it [00:00, 69.73it/s]
Validation Loss : 0.2826997827399861
100%|██████████| 170/170 [00:06<00:00, 25.41it/s]
Training Loss : 0.001024005887687535
7it [00:00, 63.19it/s]
Validation Loss : 0.28204734217036853
100%|██████████| 170/170 [00:06<00:00, 24.97it/s]
Training Loss : 0.0006642886707970106
22it [00:00, 69.32it/s]
Validation Loss : 0.28104662827470084
94%|██████████| 159/170 [00:06<00:00, 26.00it/s]
Training Loss : 0.0004568449993230923
```



We will do the same for the other dataset i.e. Multi NLI

```
multi_nli_elmo_validation = ElmoDataset(multi_nli_validation)

trainer2 = ElmoTrainer(epochs=20,lr=0.001,batch_size=64,print_every=1,device='cuda')

elmo2 = ELMo(len(multi_nli_elmo_train.vocab),100,100,80,embedding_matix)

trainer2.train(elmo2,'model/elmo2.pt',multi_nli_elmo_train,multi_nli_elmo_validation)
```

```
100% ██████████ 625/625 [02:05<00:00, 4.99it/s]
Training Loss : 0.22713012924194337
250it [00:24, 10.15it/s]
Validation Loss : 0.1599109906256199
100% ██████████ 625/625 [02:06<00:00, 4.95it/s]
Training Loss : 0.12823632835149765
250it [00:24, 10.23it/s]
Validation Loss : 0.08191098119318485
```

Sentiment Analysis Task

We use the elmo embeddings and the corresponding lstm layers in our Sentiment Classification model and train it after fiddling with the hyper-parameters.

```
class SentimentClassifier(nn.Module):
    def __init__(self, embedding_dim, hidden_dim, elmo_embeddings, elmo_l1, elmo_l2, dropout=0.2):
        super(SentimentClassifier, self).__init__()
        self.embedding_dim = embedding_dim
        self.hidden_dim = hidden_dim
        self.embedding = nn.Embedding.from_pretrained(elmo_embeddings)
        self.embedding.weight = nn.Parameter(self.embedding.weight, requires_grad=True)
        self.weights = nn.Parameter(torch.tensor([0.33, 0.33, 0.33]), requires_grad=True)
        self.linear1 = nn.Linear(embedding_dim, hidden_dim*2)
        self.lstm1 = elmo_l1
        self.lstm2 = elmo_l2
        self.linear2 = nn.Linear(hidden_dim*2, 1)
        self.dropout = nn.Dropout(dropout)

    def forward(self, input_data):
        embeds = self.embedding(input_data)
        embeds_change = self.linear1(embeds)
        hidden1, _ = self.lstm1(embeds)
        hidden2, _ = self.lstm2(hidden1)
        elmo_embed = (self.weights[0]*hidden1 + self.weights[1]*hidden2
                     + self.weights[2]*embeds_change)/(self.weights[0]+self.weights[1]+self.weights[2])
        elmo_embed_max = torch.max(elmo_embed, dim=1)[0]
        elmo_embed_max_drop = self.dropout(elmo_embed_max)
        linear_out = self.linear2(elmo_embed_max_drop)
        return torch.sigmoid(linear_out)]
```





```

Training Accuracy : 0.8399999999999999
100% ██████████ 22/22 [00:00<00:00, 136.53it/s]
Validation Accuracy : 0.48454545454545456
100% ██████████ 170/170 [00:02<00:00, 64.61it/s]
Training Accuracy : 0.8610588235294121
100% ██████████ 22/22 [00:00<00:00, 125.75it/s]
Validation Accuracy : 0.48363636363636364
100% ██████████ 170/170 [00:02<00:00, 62.43it/s]
Training Accuracy : 0.888352941176471
100% ██████████ 22/22 [00:00<00:00, 134.65it/s]
Validation Accuracy : 0.48727272727272725
100% ██████████ 170/170 [00:02<00:00, 64.19it/s]
Training Accuracy : 0.9108235294117644
100% ██████████ 22/22 [00:00<00:00, 136.98it/s]
Validation Accuracy : 0.49272727272727274
100% ██████████ 170/170 [00:02<00:00, 63.65it/s]
Training Accuracy : 0.9277647058823527
100% ██████████ 22/22 [00:00<00:00, 137.59it/s]
Validation Accuracy : 0.4890909090909091

```

We plotted the training losses and accuracy. Classification report not possible as initially it was a regression task so instead approximated using a threshold.





Natural Language Inference

For this task I have just created the Elmo Embeddings not created any model.

```
class MultiNliData(DataPipeline):
    def __init__(self, filename, type, max_seq_len=50, min_freq=3, sentence_limit=50000, vocab=None):
        self.sentence_limit = sentence_limit
        super().__init__(filename, type, max_seq_len, min_freq, vocab)
    def read_data(self, filename, type):
        datacleaner = DataCleaner()
        data = load_from_disk(filename)
        processed_data = []
        premise = []
        hypothesis = []
        target = []
        counter = 0
        for line in tqdm(data[type]):
            if counter > self.sentence_limit:
                break
            counter += 1
            p = datacleaner.process(line['premise']).split(" ")
            h = datacleaner.process(line['hypothesis']).split(" ")
            processed_data.append(p)
            processed_data.append(h)
            premise.append(p)
            hypothesis.append(h)
            target.append(line['label'])
        self.data = processed_data
        self.target = target
        self.premise = premise
```

```
def mnli_training():
    multi_nli_train = MultiNliData('data/multi_nli.hf', 'train', 200, 3, 20000)
    multi_nli_validation = MultiNliData('data/multi_nli.hf', 'validation_matched', 200, 3, 8000, multi_nli_train.get_vocab())
    multi_nli_elmo_train = ElmoDataset(multi_nli_train)
    multi_nli_elmo_validation = ElmoDataset(multi_nli_validation)
    glove = load_embeddings(multi_nli_elmo_train.vocab, "data/glove.6B.100d.txt", 100)
    trainer2 = ElmoTrainer(epochs=20, lr=0.001, batch_size=64, print_every=1, device='cuda')
    elmo2 = ELMo(len(multi_nli_elmo_train.vocab), 100, 100, 200, glove)
    trainer2.train(elmo2, 'model/elmo2.pt', multi_nli_elmo_train, multi_nli_elmo_validation)
    trainer2.plot_loss()
```



Conclusion

This assignment taught us a lot.

